

Bringing Riak to the Mobile Platform

Kresten Krab Thorup
Hacker

@drkrab



TRIFORK.

About the Speaker

- **Language Geek** Emacs/TeX Hacker,
Objective C, NeXT, GNU Compiled Java,
Java Generics, Ph.D., Erlang/Erjang
- **Developer** J2EE AppServer,
CORBA/RMI, XA-TM
- **Trifork CTO** QCon & GOTO Conferences
Technology Adoption; Riak; Erlang

Outline

- Riak and Protocol
- BucketDB, a Riak clients for mobile



Outline

- **Riak** and it's Data Model
- **Protocol** for Key/Value synchronization
- **BucketDB**, a Riak clients for mobile

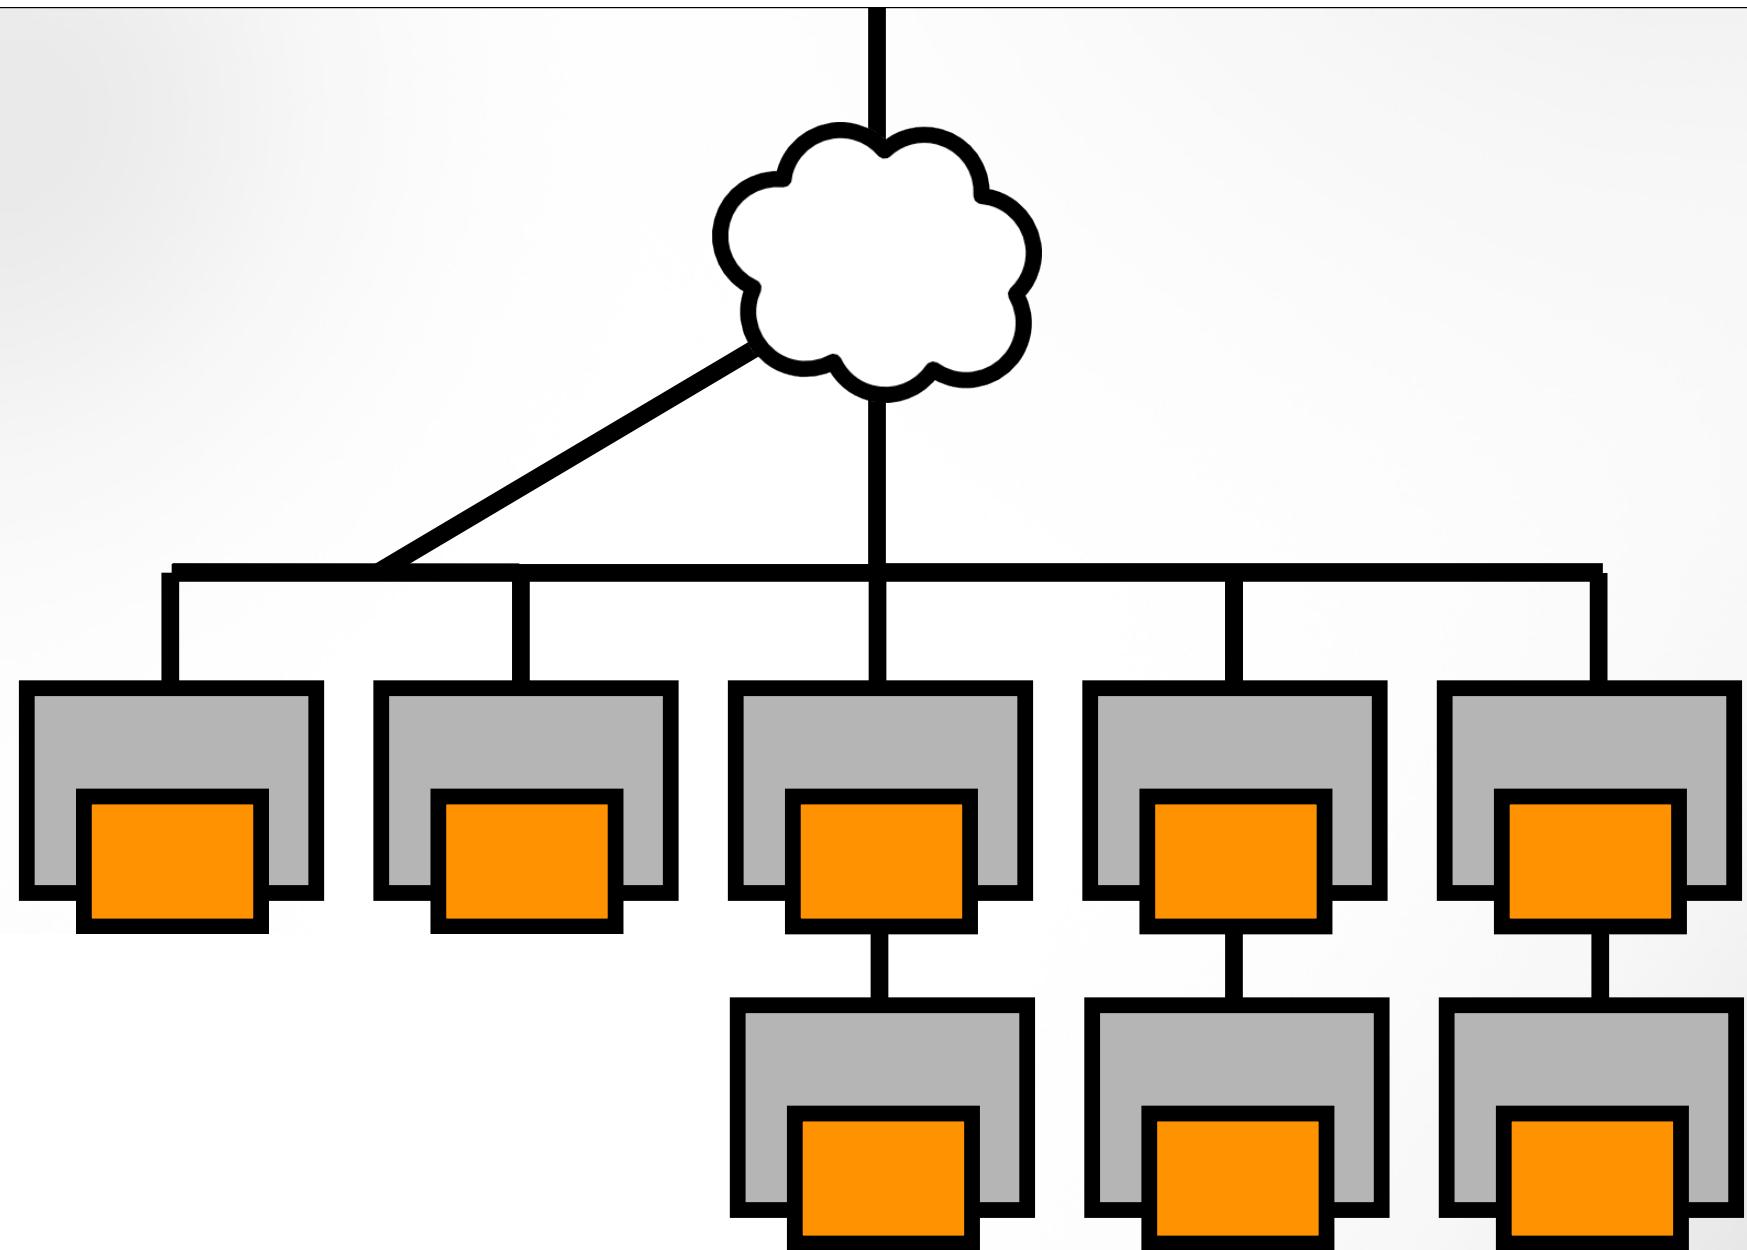
@drkrab

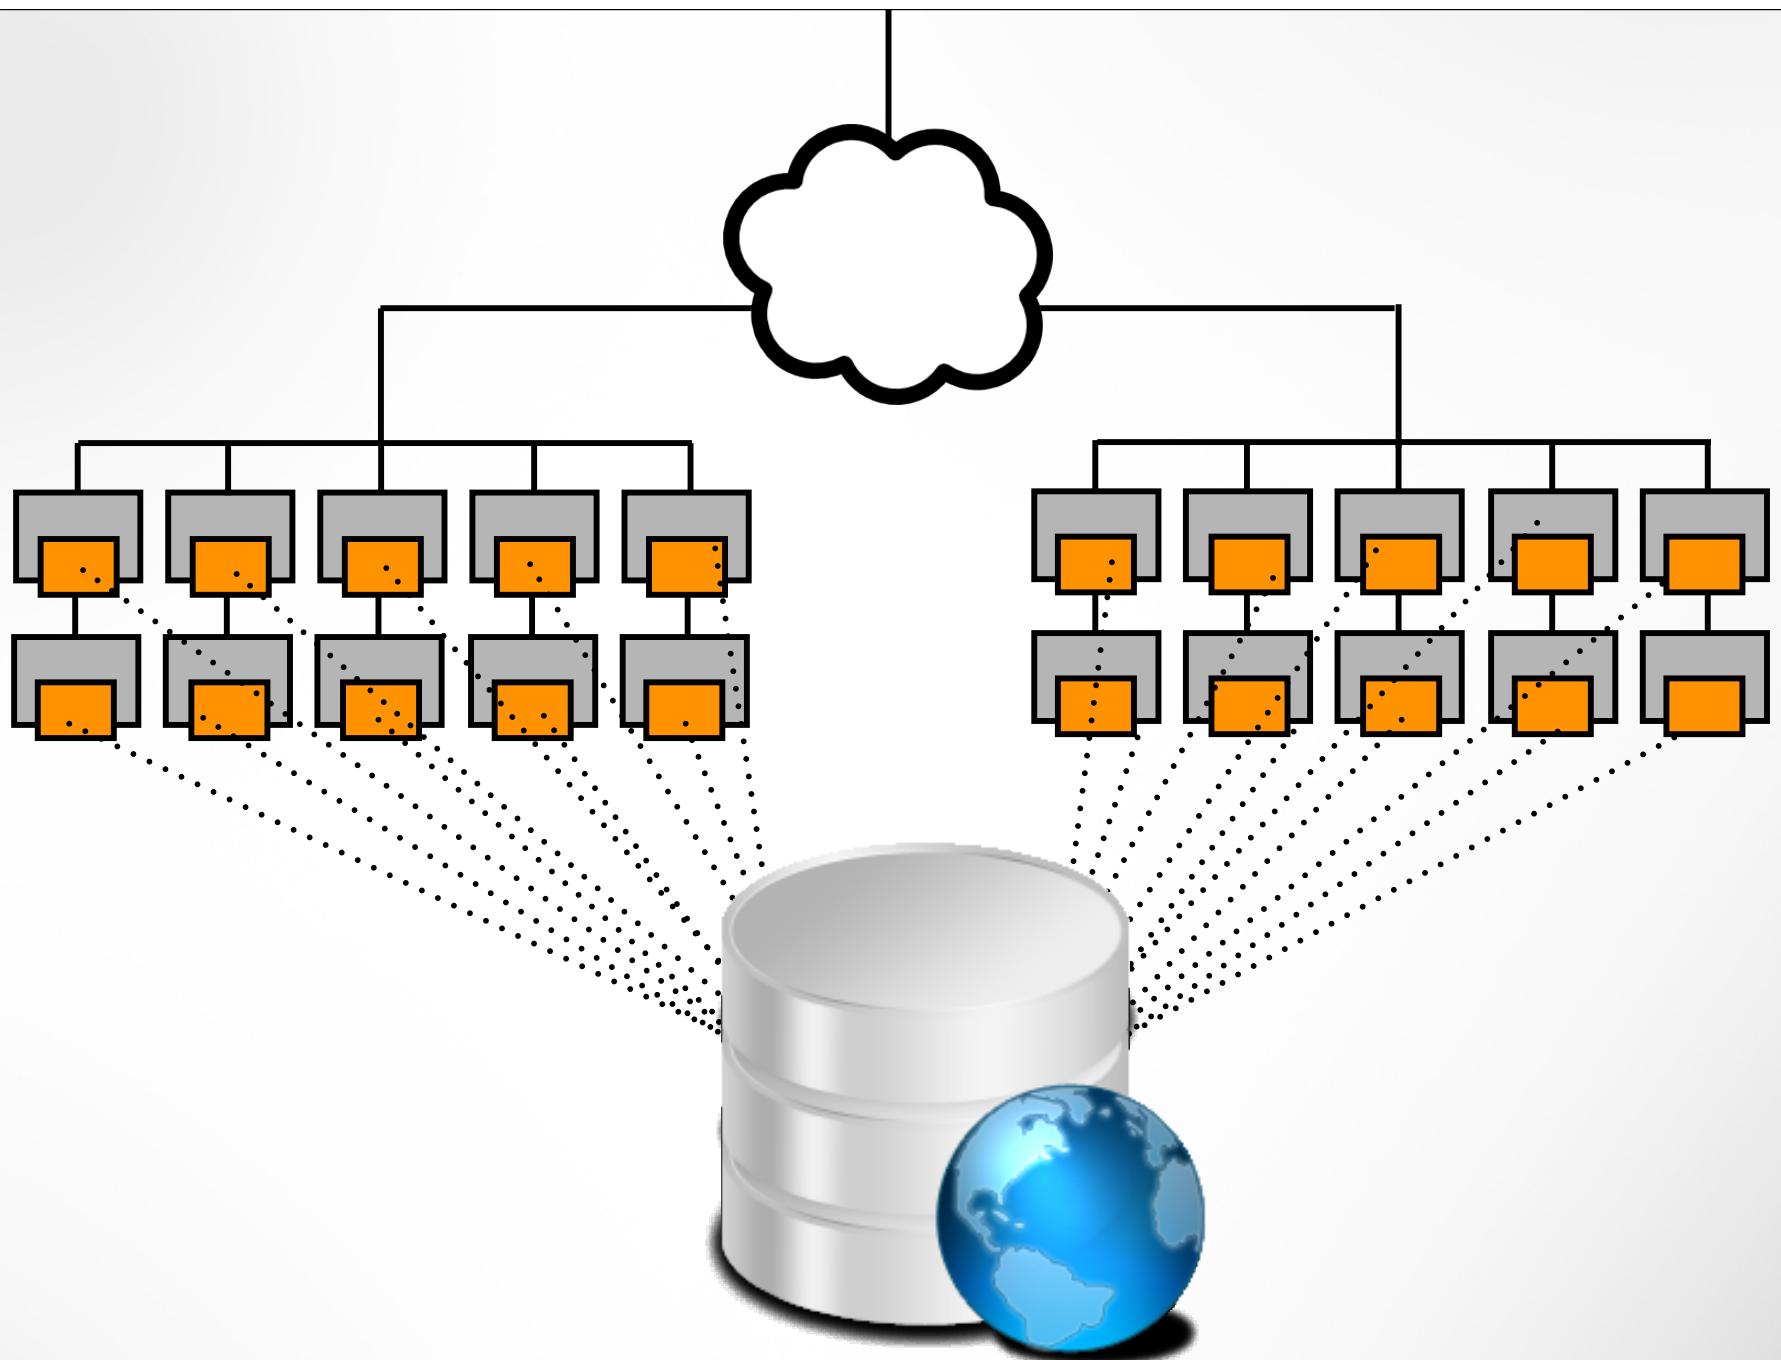


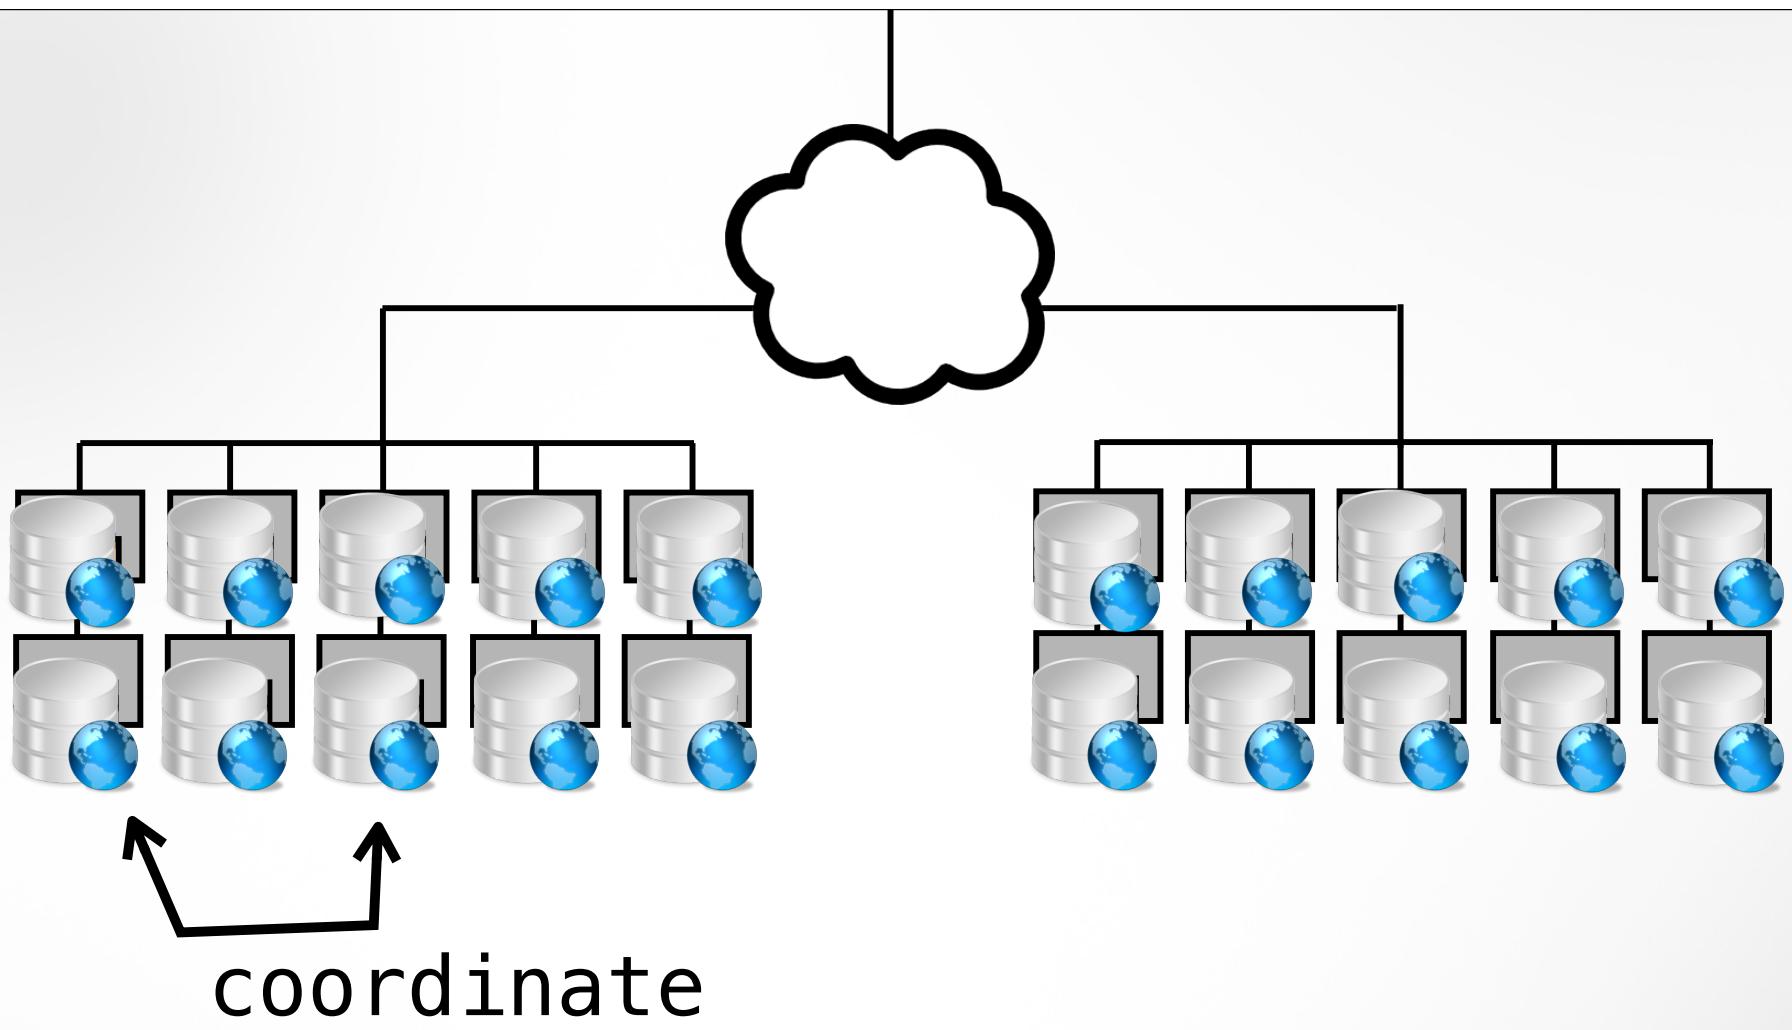
What is Riak, anyway?

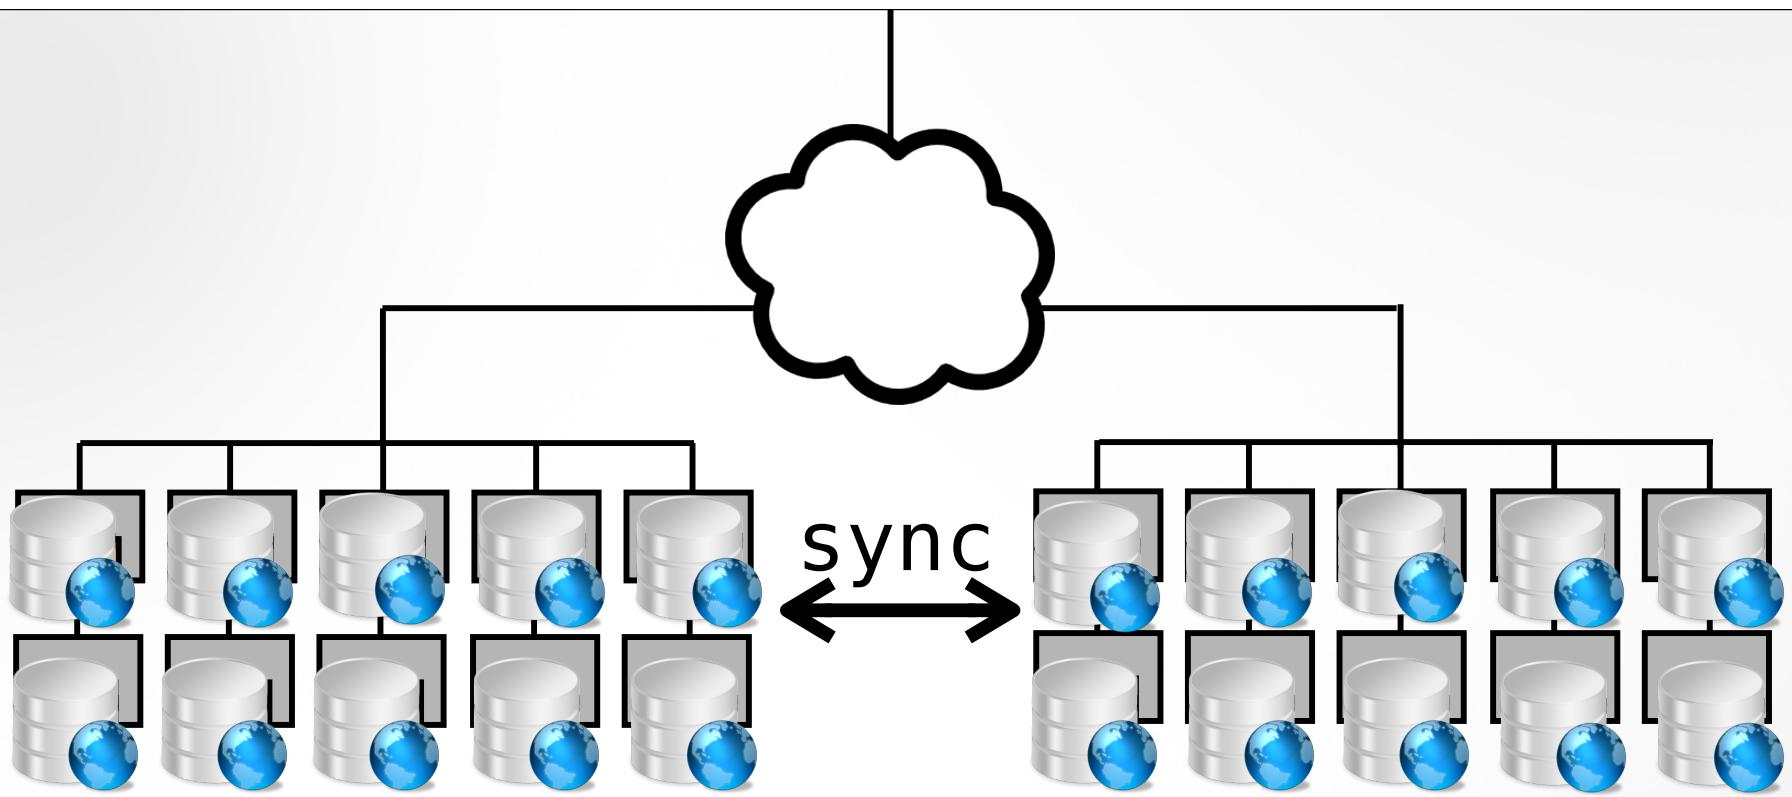
@drkrab

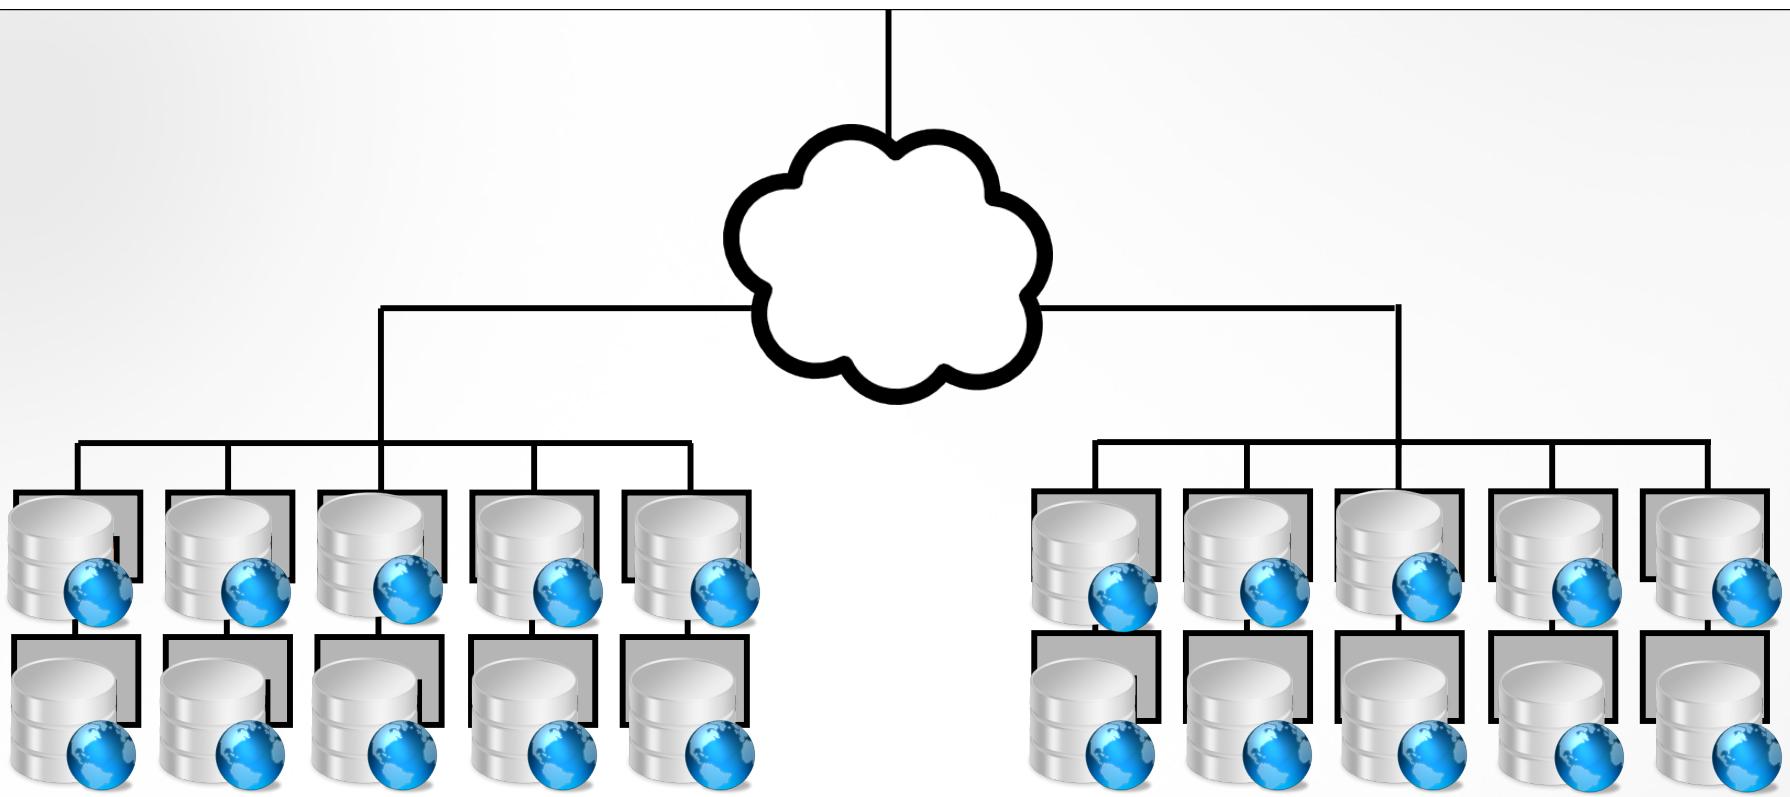








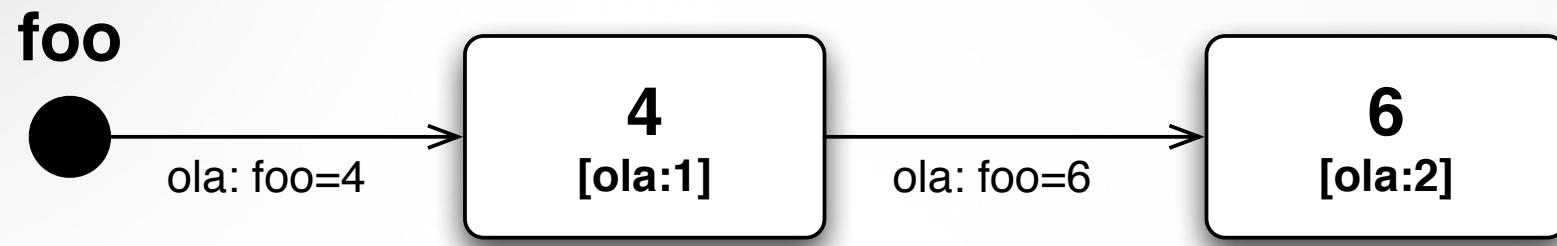


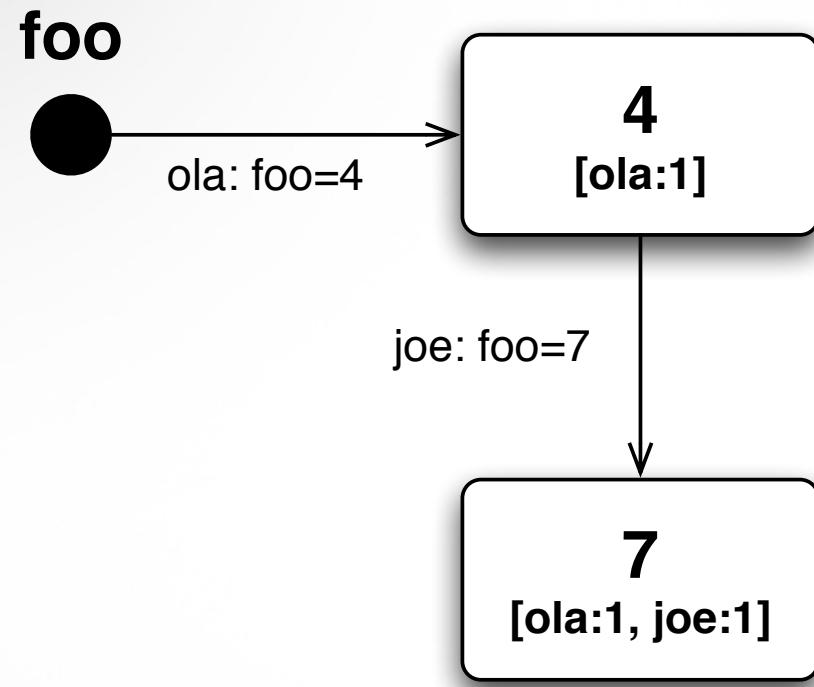


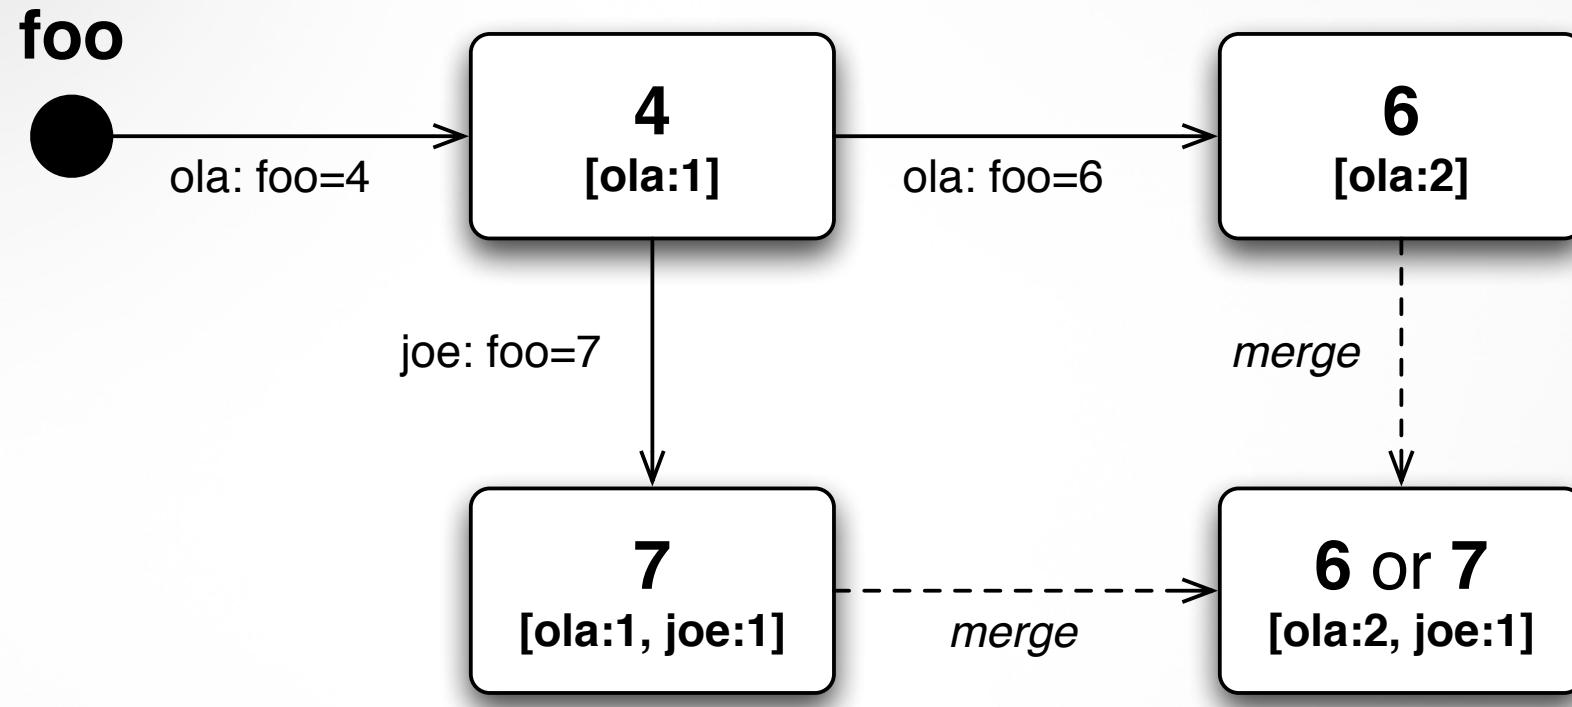
- scalable and available
- system captures write conflicts
- resolve lazily (read repair)

Shared Medicine Card









Pseudo Code!

Riak API review

connect(ClientID) -> Client

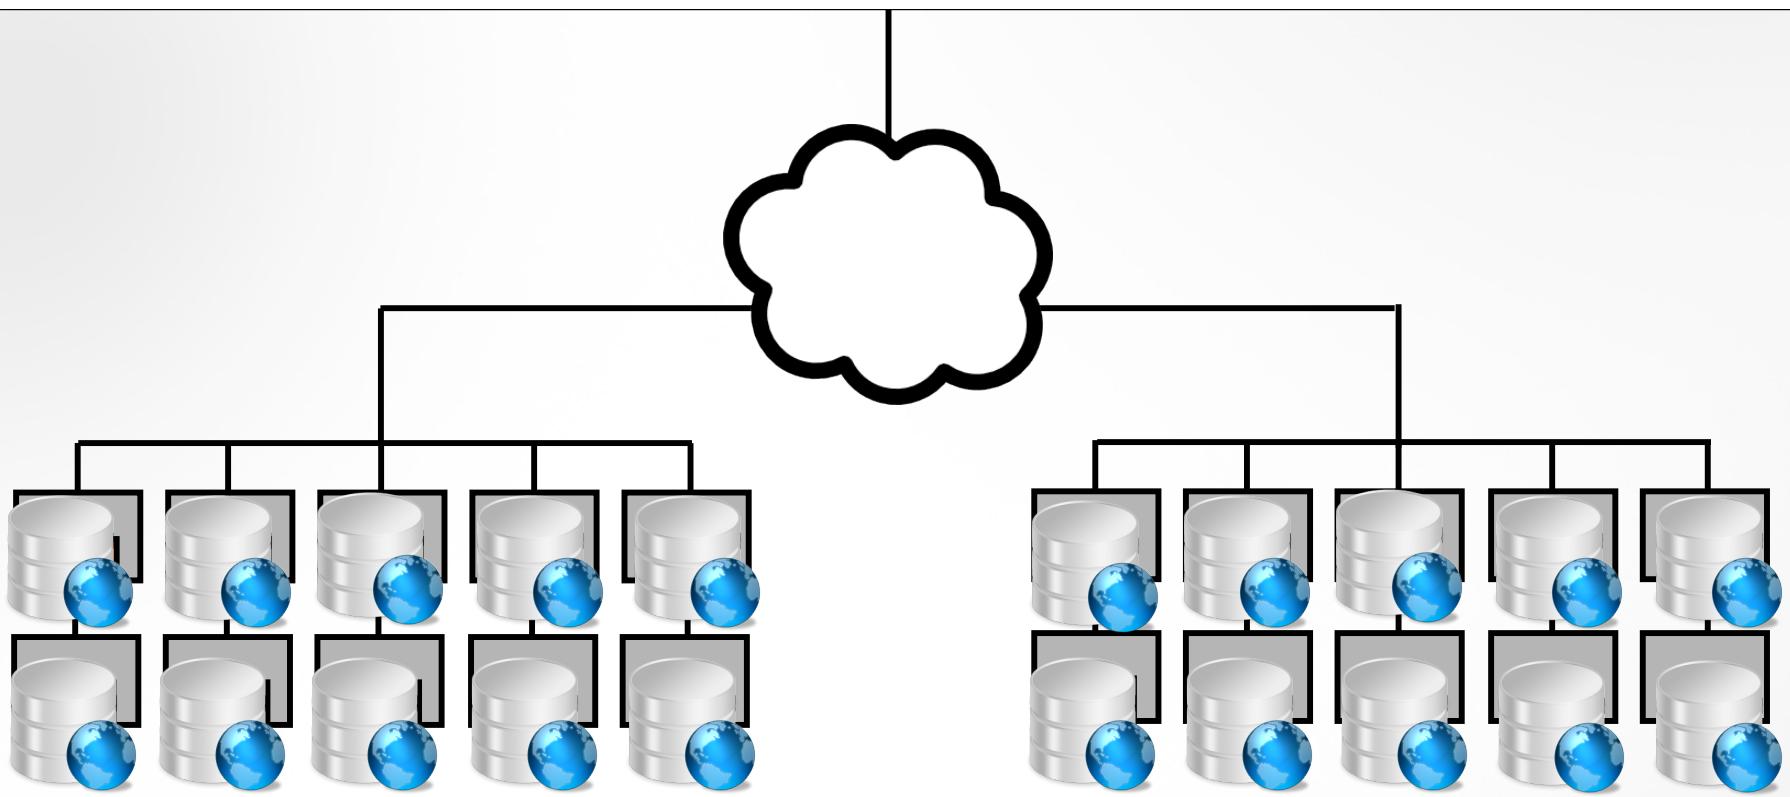
```
BKey :: {Bucket, Key}  
Datum :: {ContentType, RawData}
```

Client:**insert(BKey, Datum)** -> {ok, VClock} | ...

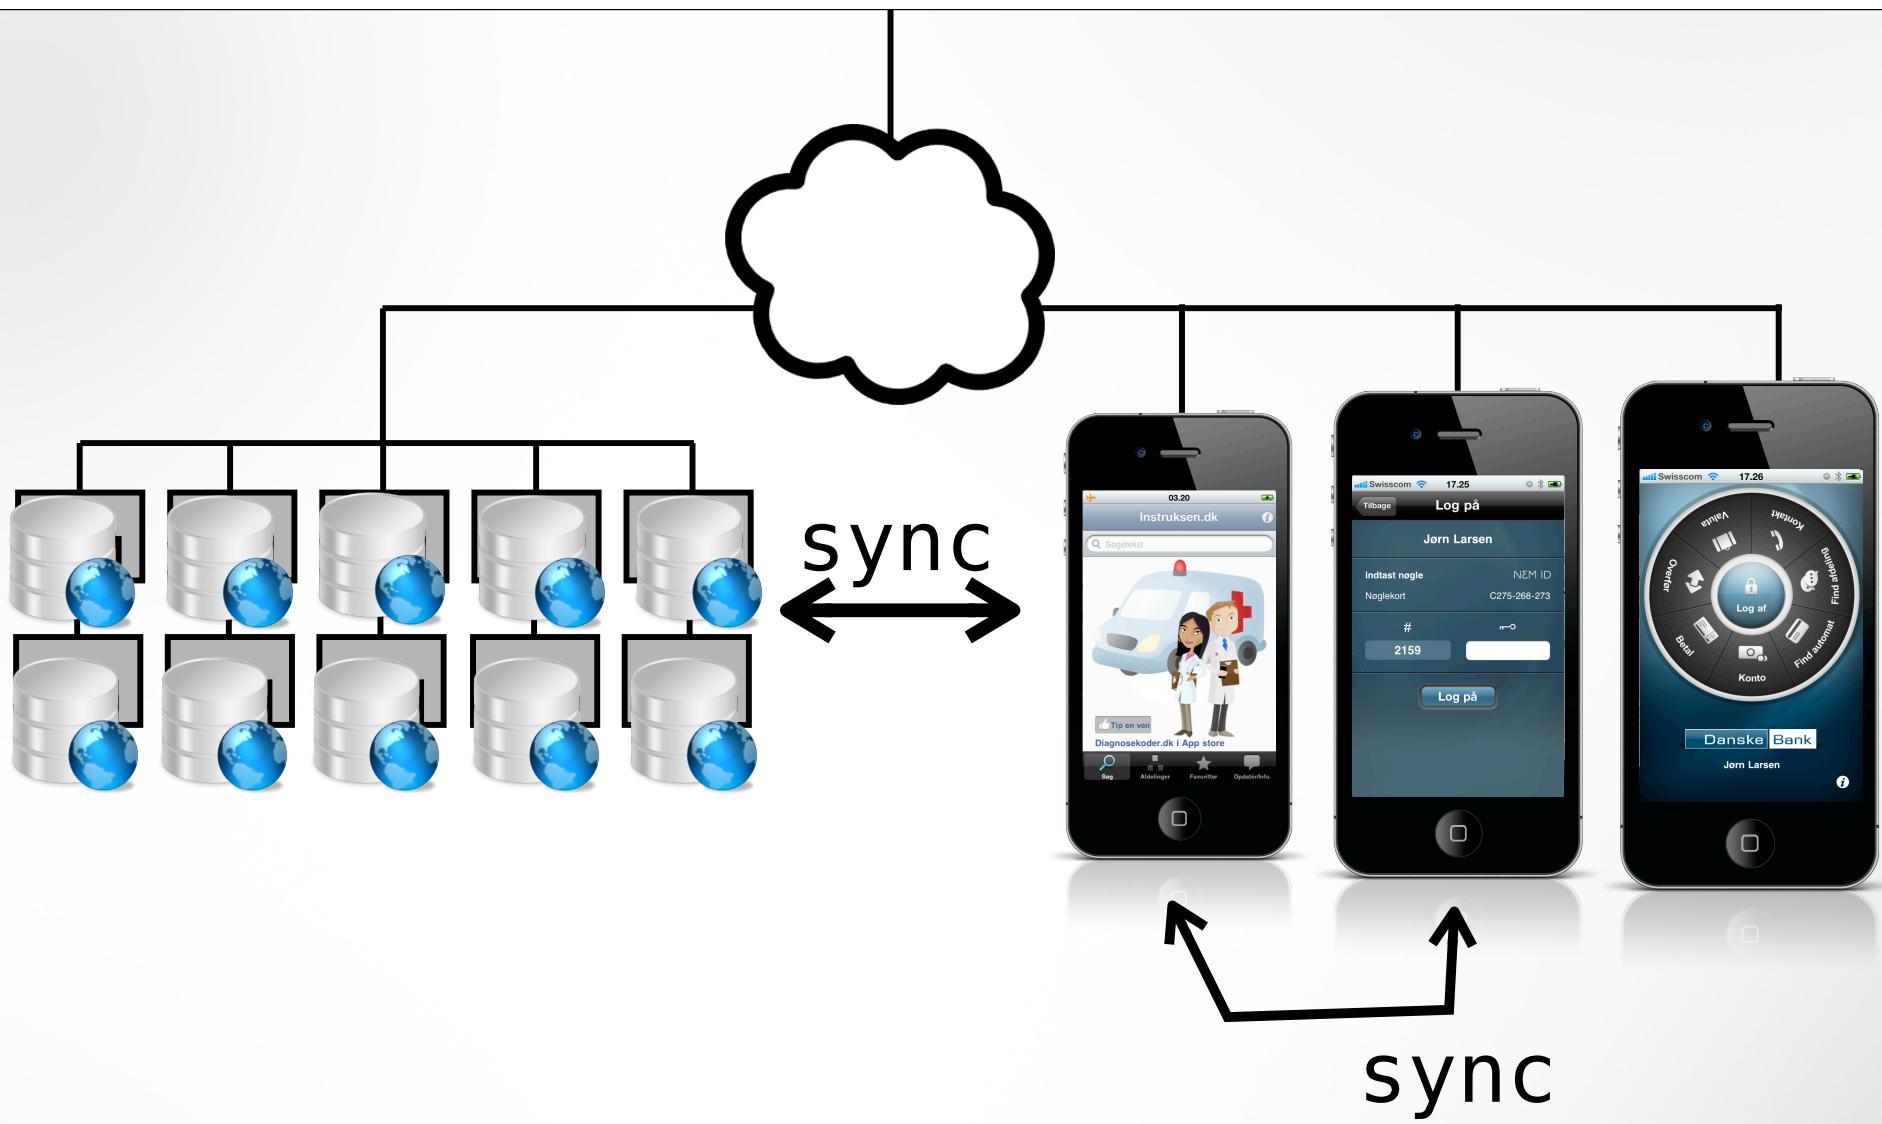
Client:**lookup(BKey)** -> {ok, VClock, [Datum]} | ...

Client:**update(BKey, VClock, Datum)** -> ok | ...

Client:**delete(BKey, VClock)** -> ok | ...



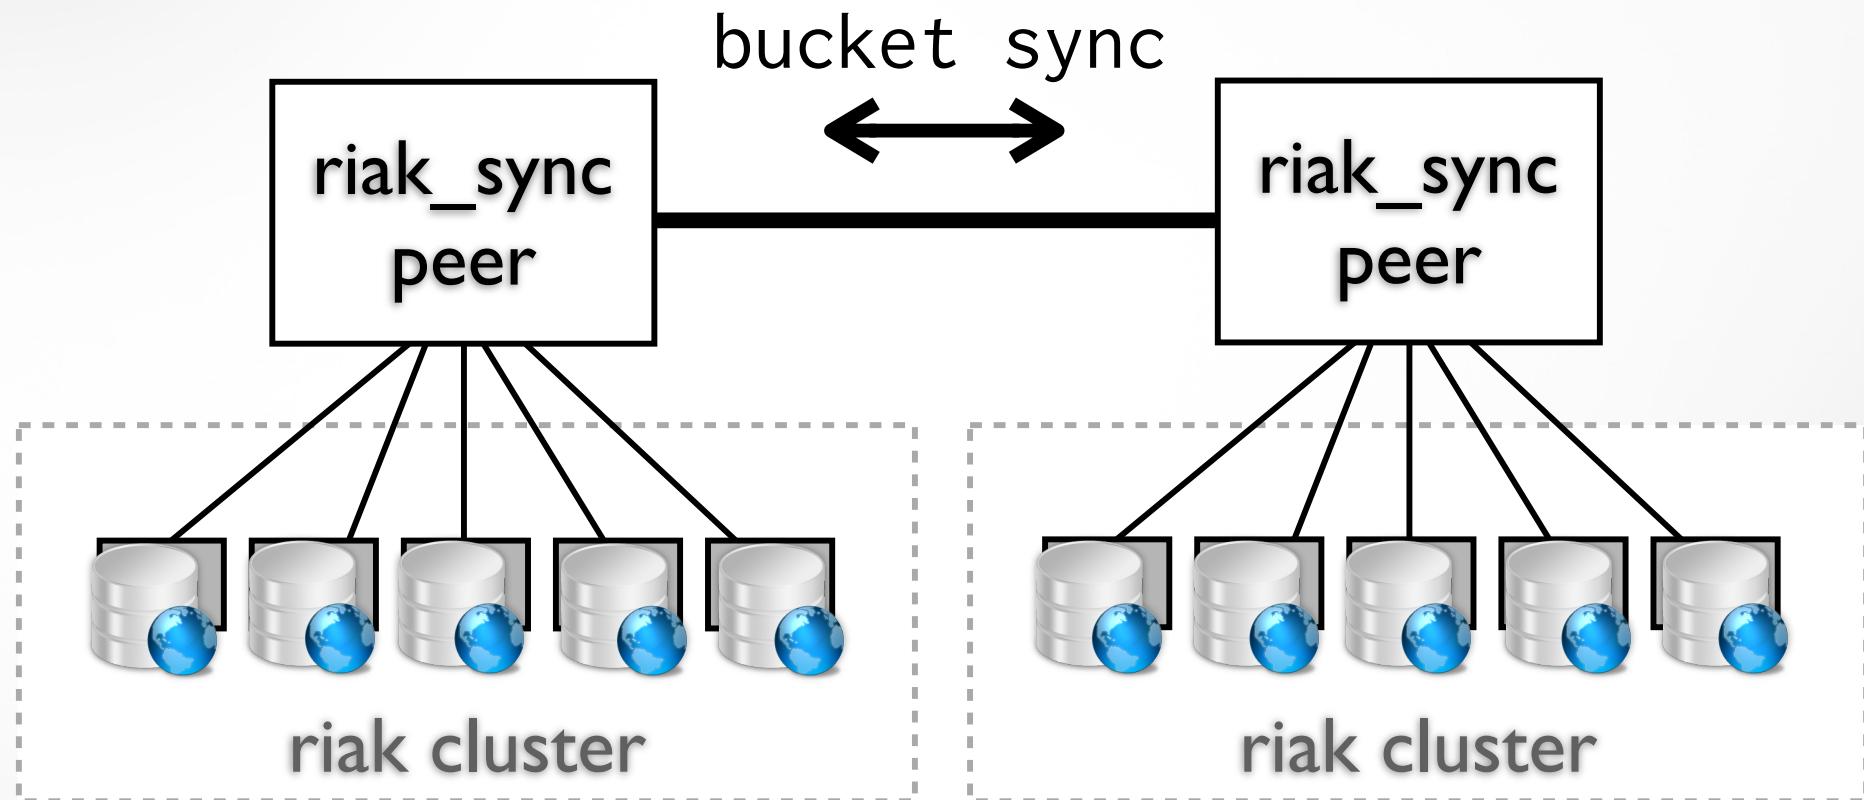
- scalable and available
- system captures write conflicts
- resolve lazily (read repair)



Today: Two Things

- **RiakSync**, Two-way Replication Protocol
- **BucketDB**, the Mobile client

Not to be confused with `riak_repl`,
part of the commercial Riak EDS



```
%% “client” peer
> LocalStore = riak:client_connect('riak@127.0.0.1')
> riak_sync:sync(LocalStore, <<“bucket”>>,
    “172.1.35.204”, 8082)
```

```
%% Other site has riak_sync running...
(riak@172.1.35.204)1> riak_sync:start()
```

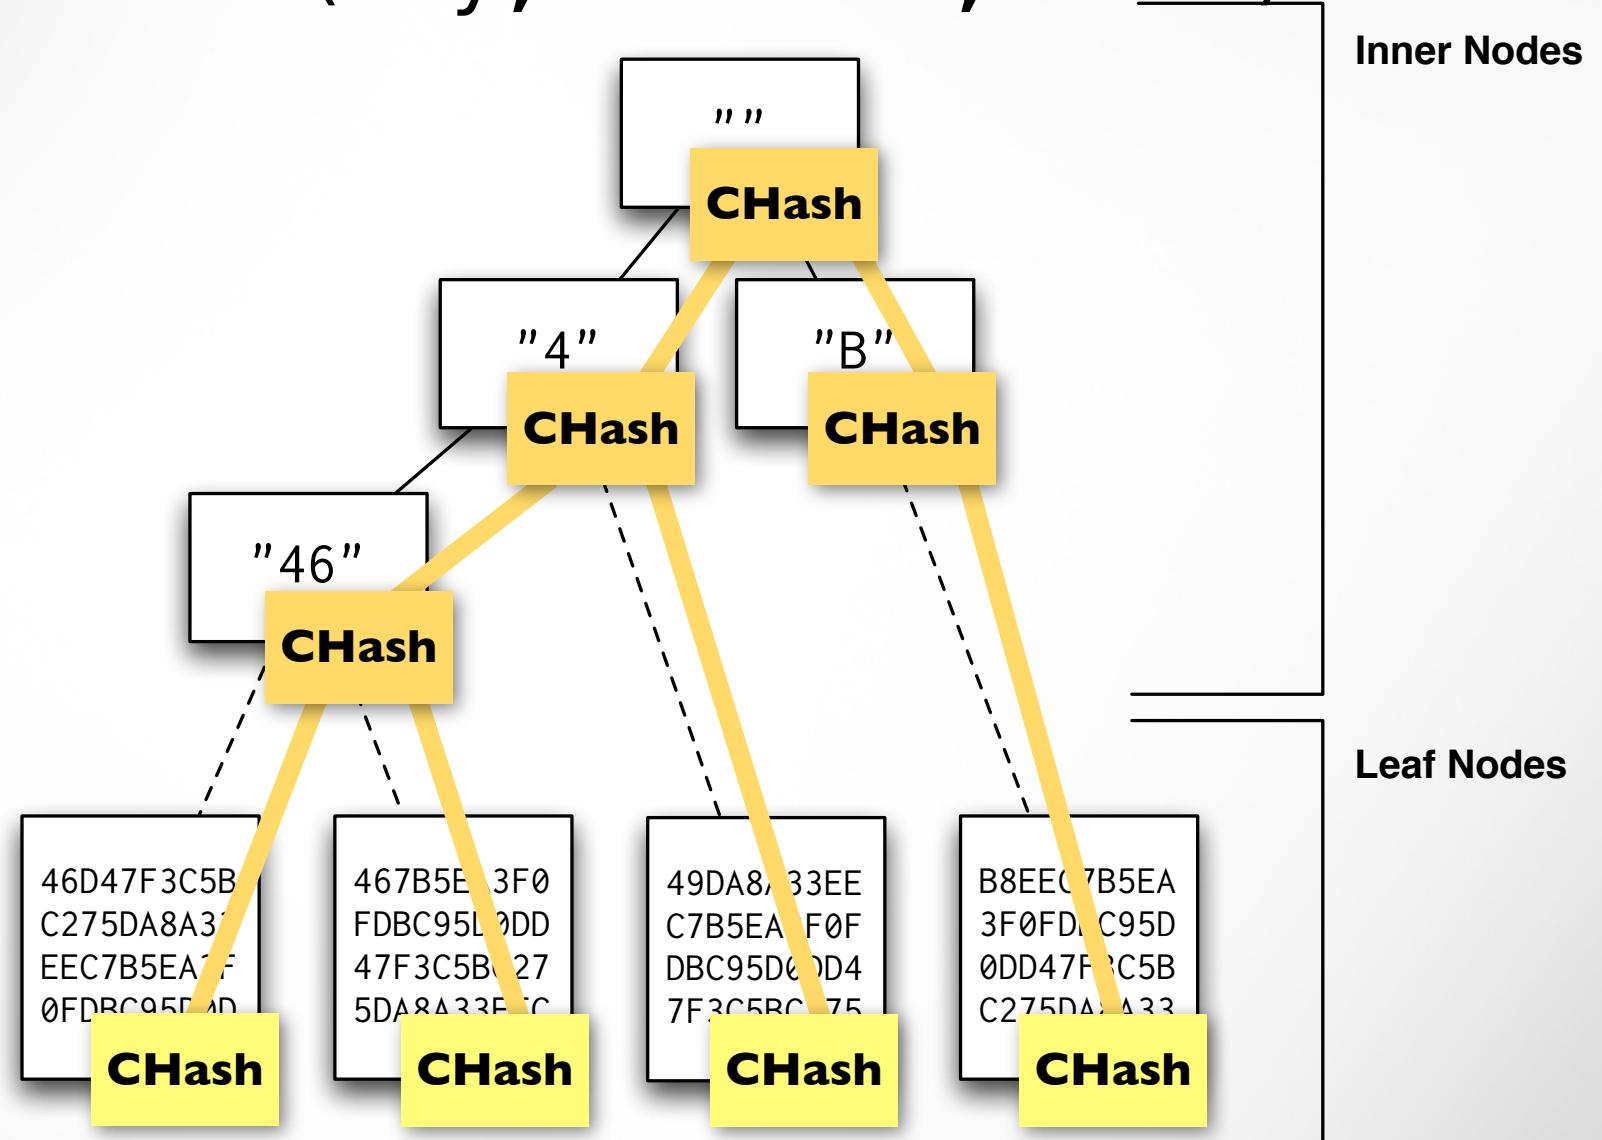
```
{riak_sync, [{pb_ip, “172.1.35.204”},
             {pb_port, 8082},
             {riak, ‘dev@127.0.0.1’}]}  
}
```

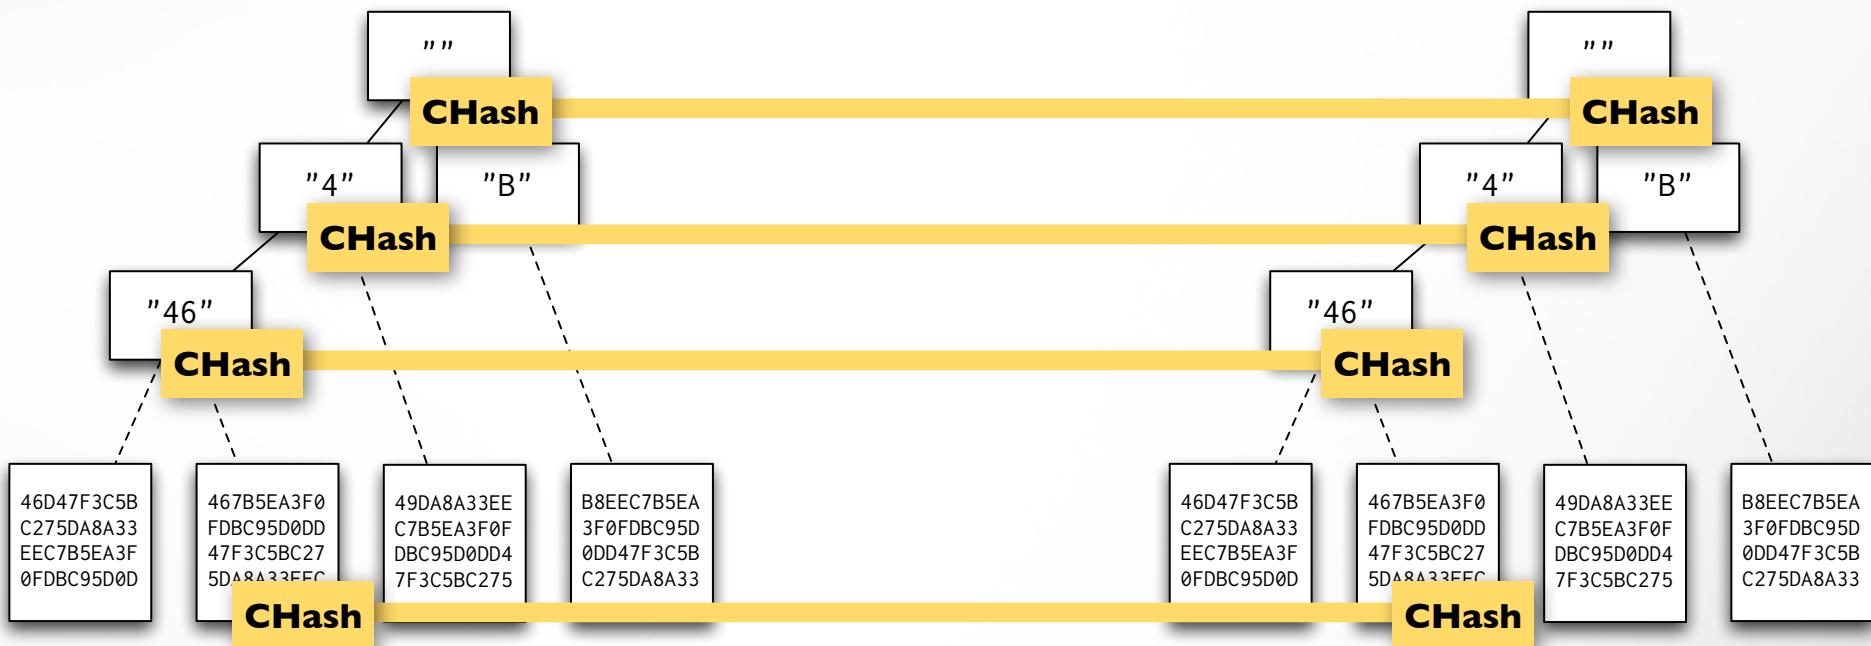
Riak Sync Protocol

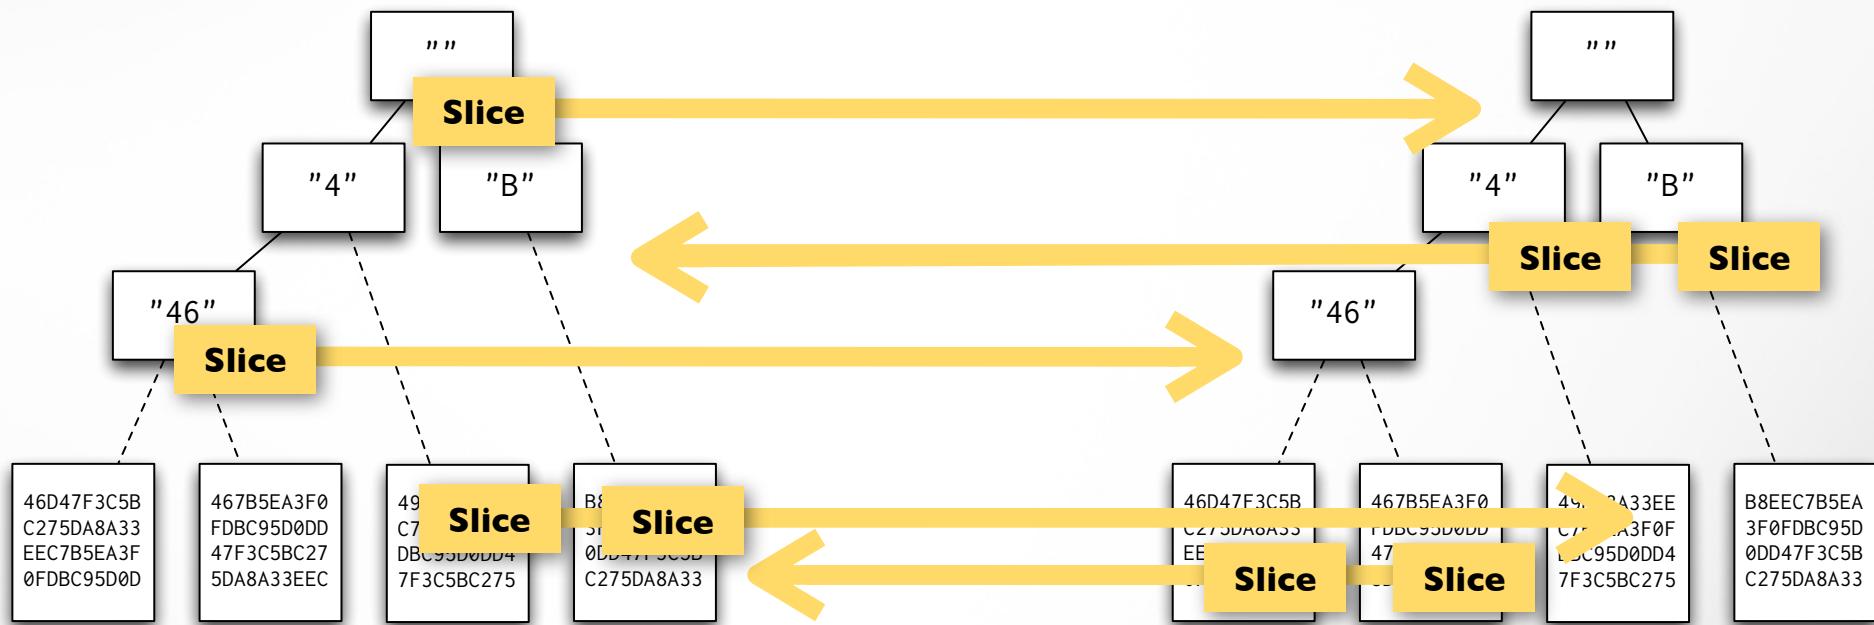
- Riak bucket replication
 - Asymmetric work load
 - Designed for high-latency networks
-
- **Very different design criteria than normal Riak cluster-cluster**

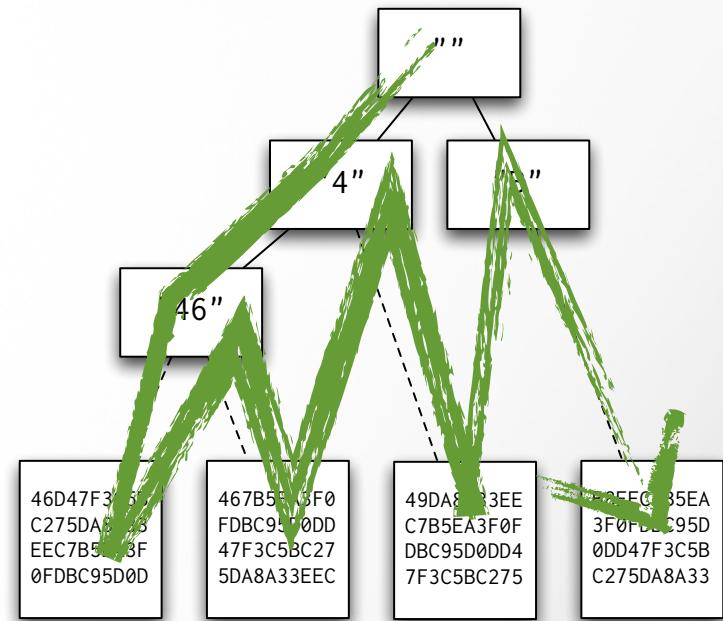
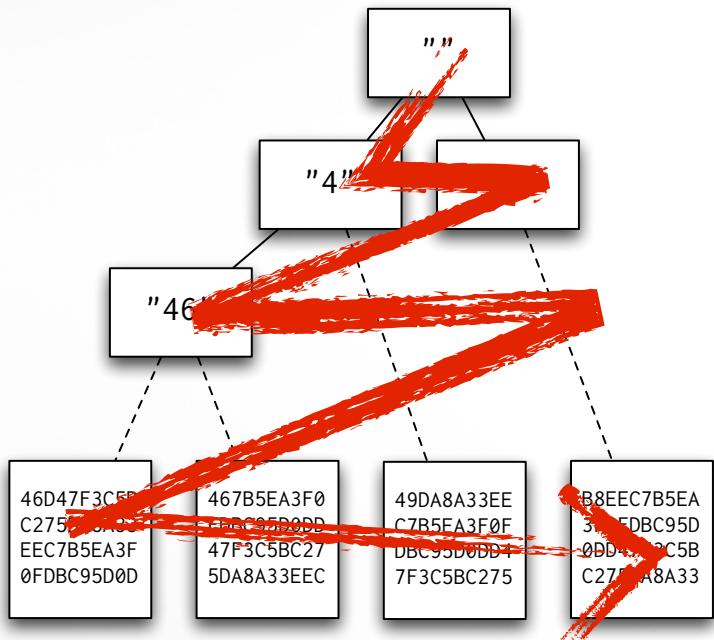
`StorageKey = sha1(Key)`

`CHash = sha1(Key, Version, Value)`



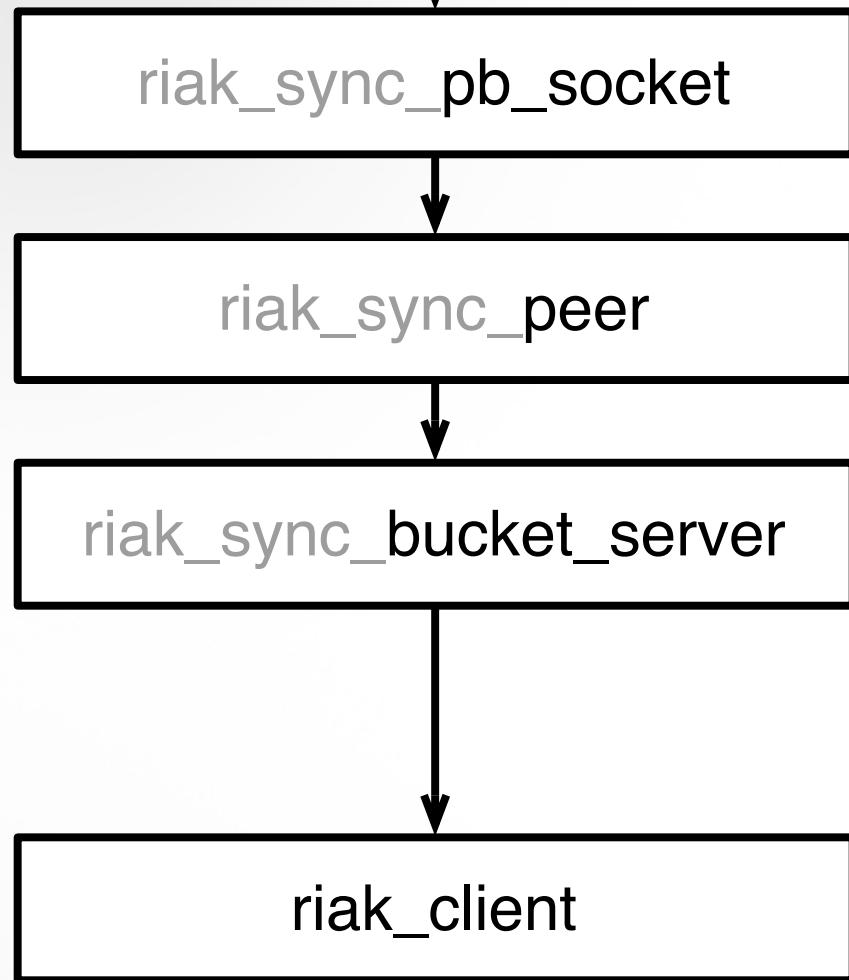






Sync Protocol

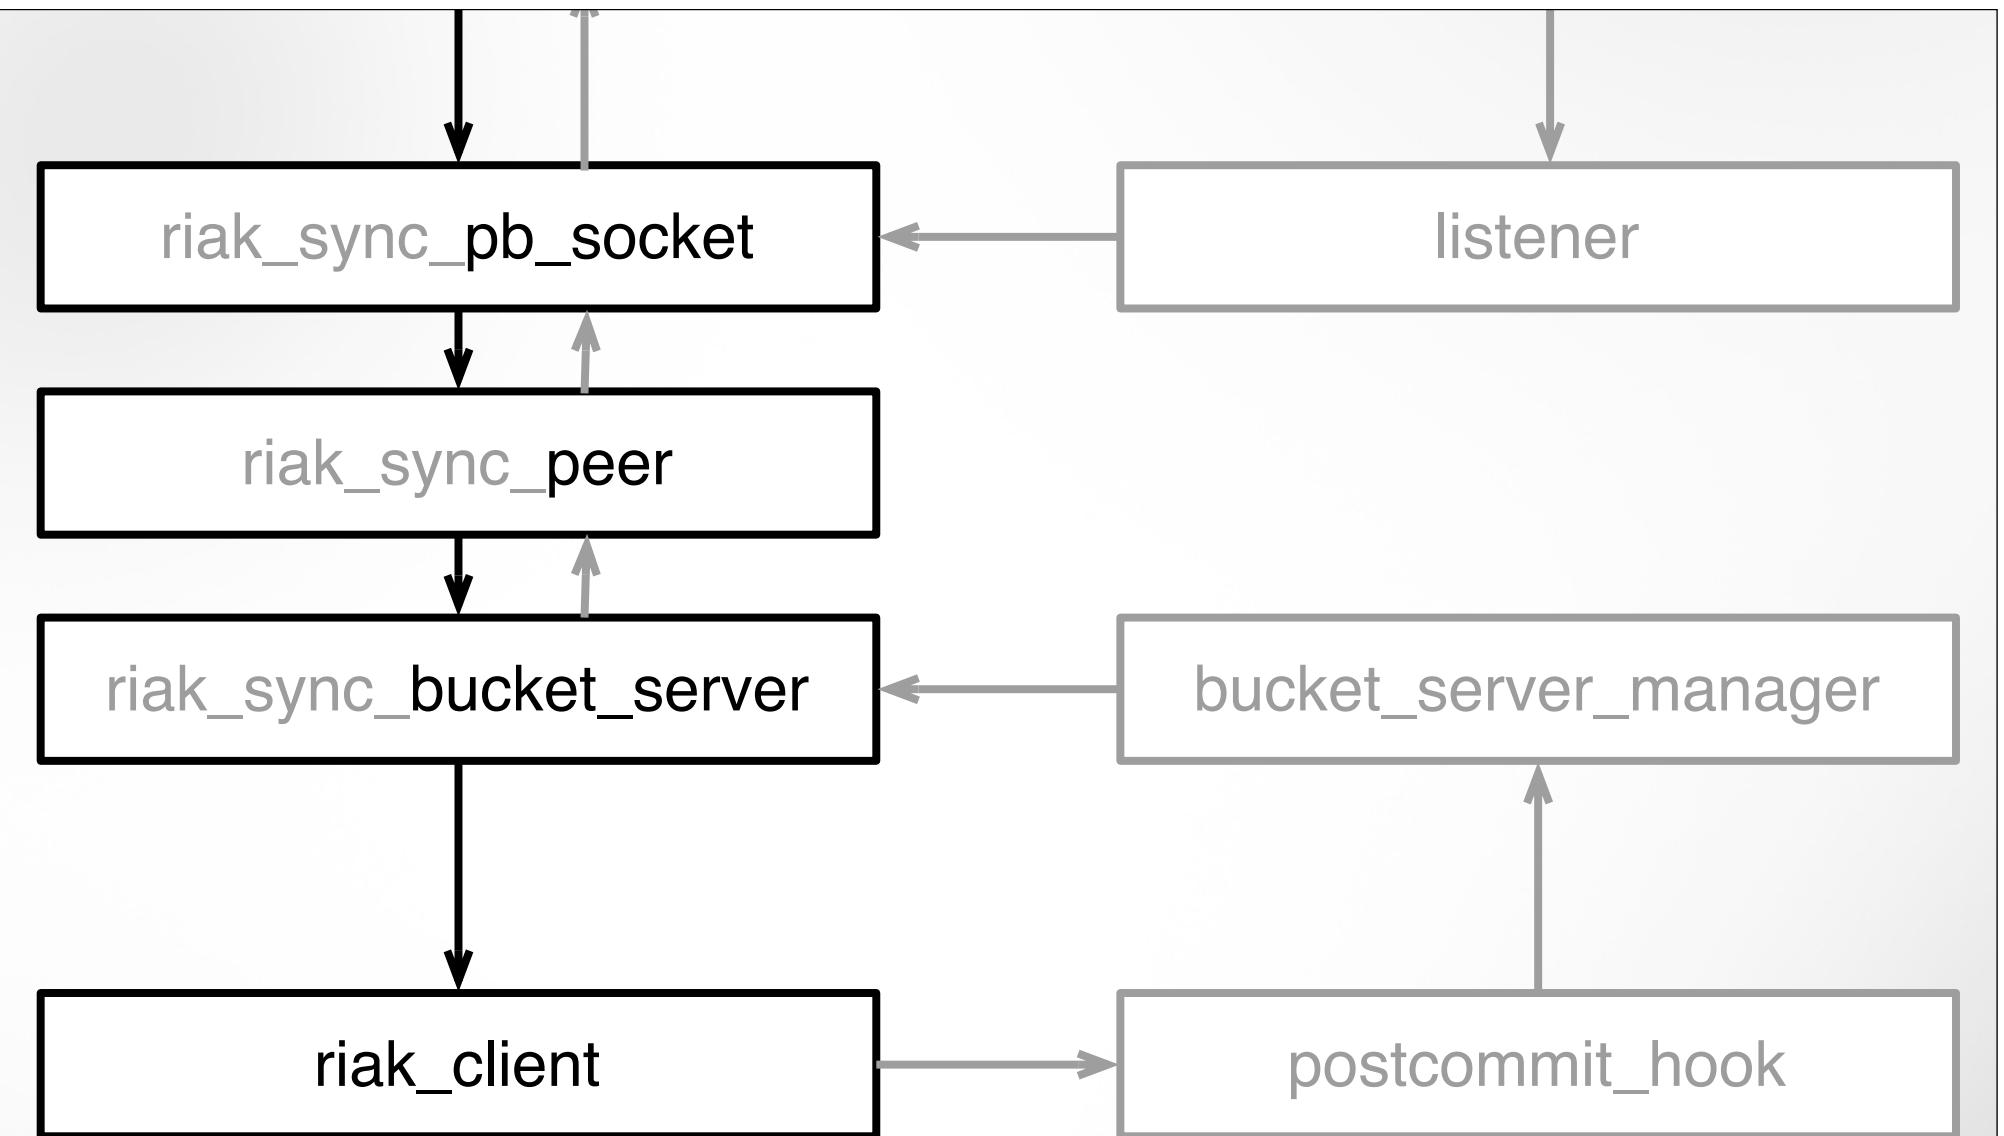
- Parallelizes quickly (client does breadth first)
- Many small messages (will benefit from Nagle's alg)
- Server peer reorders requests to reduce session/buffer state
- Client needs (almost) no session state



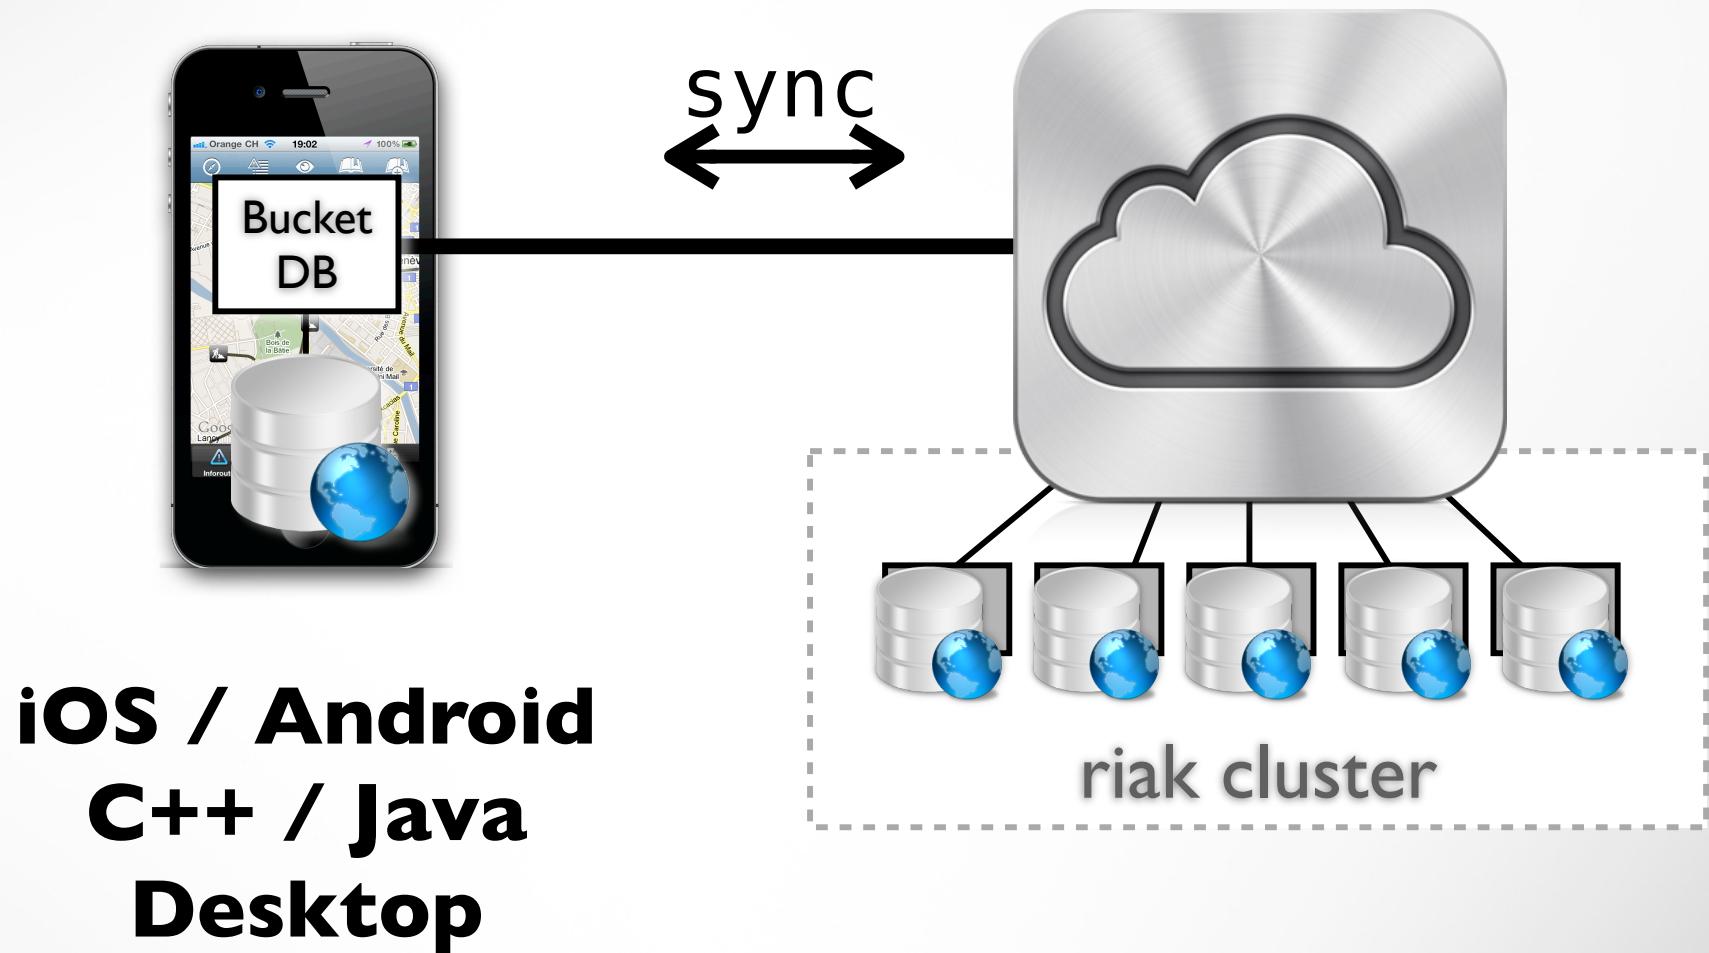
wire protocol handler

session management

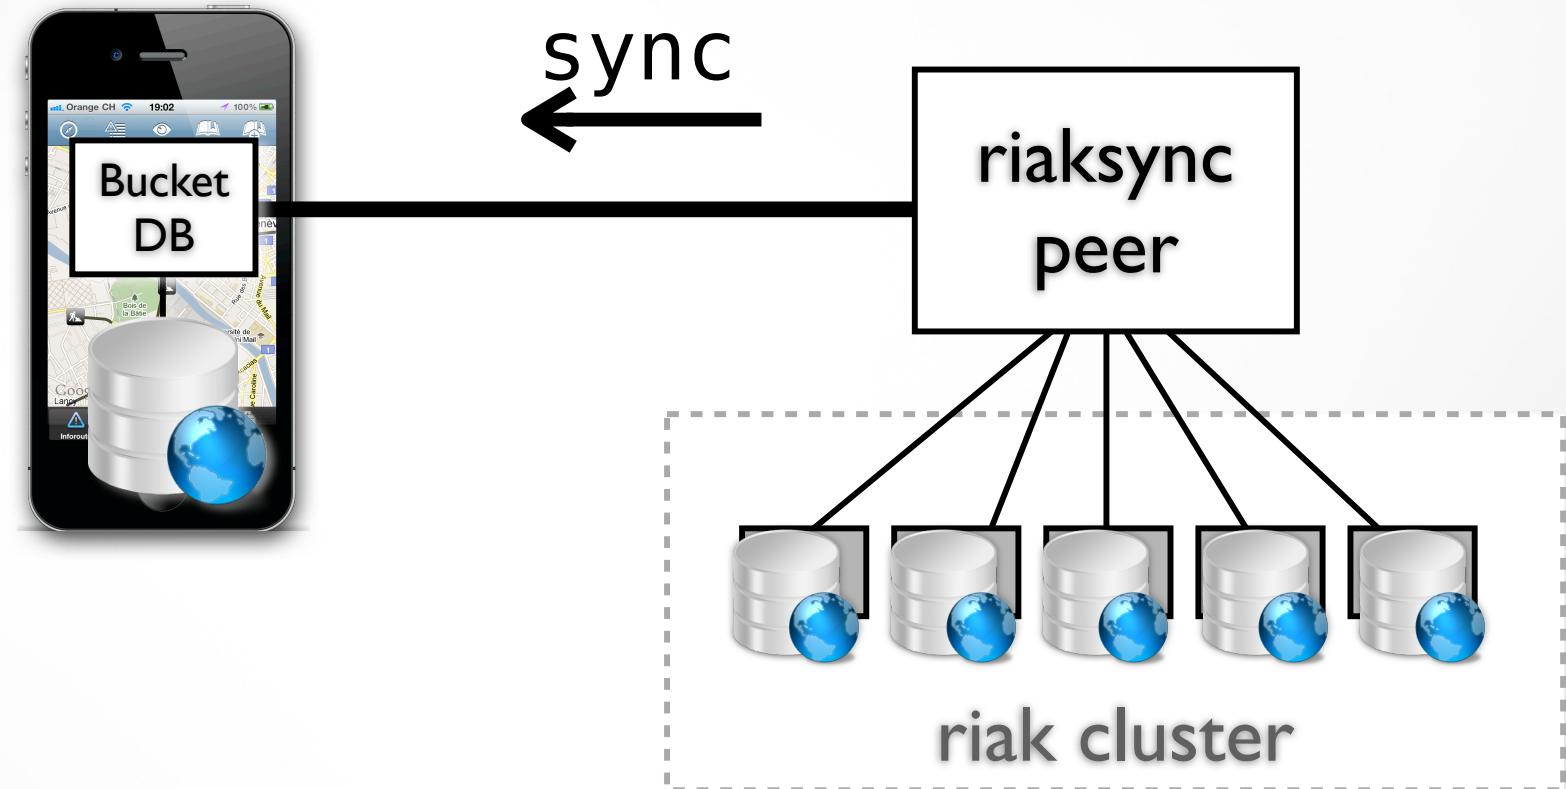
merkle_tree server



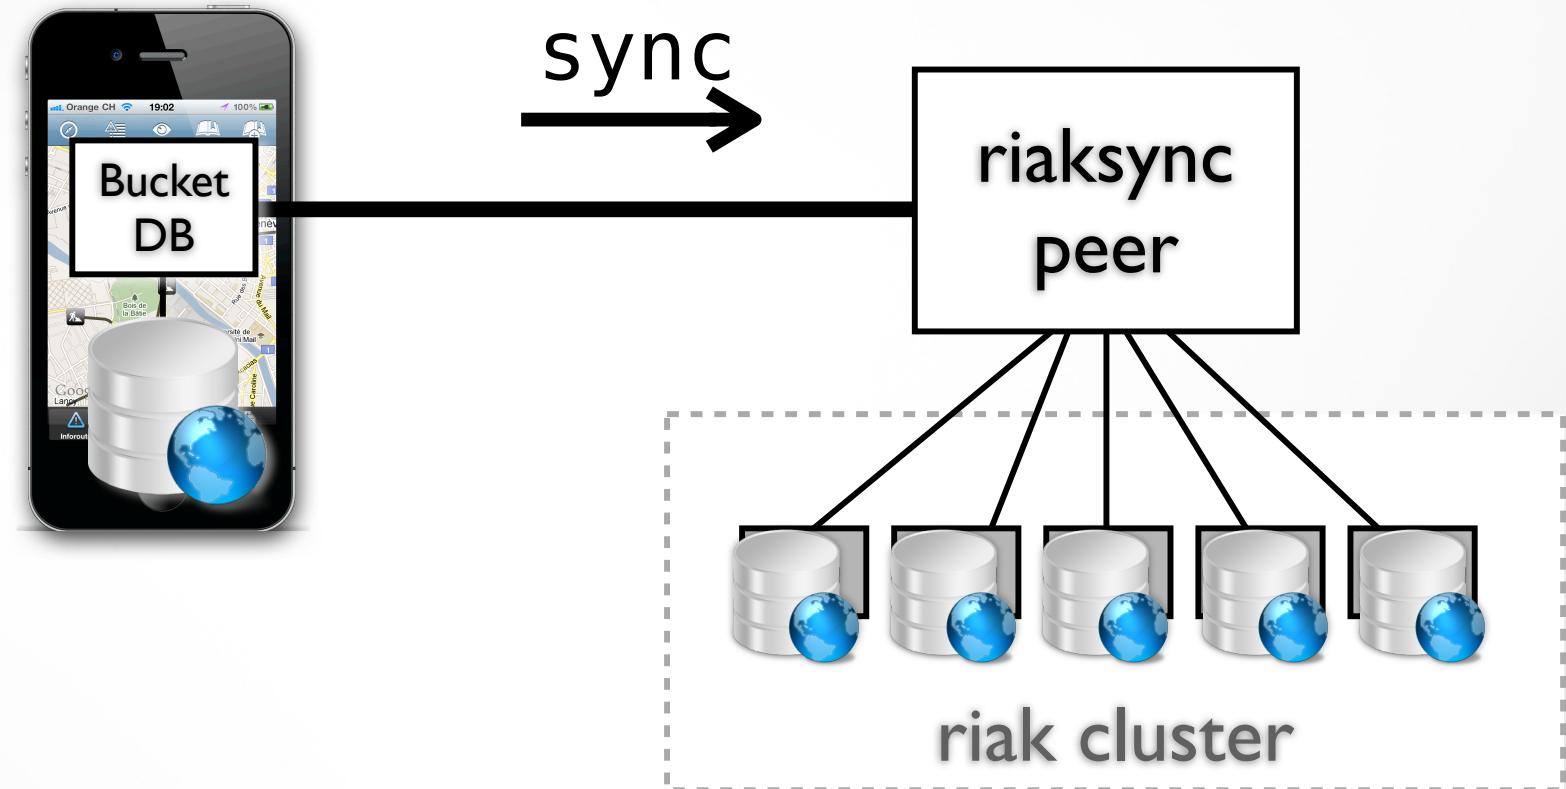
Mobile Riak: BucketDB



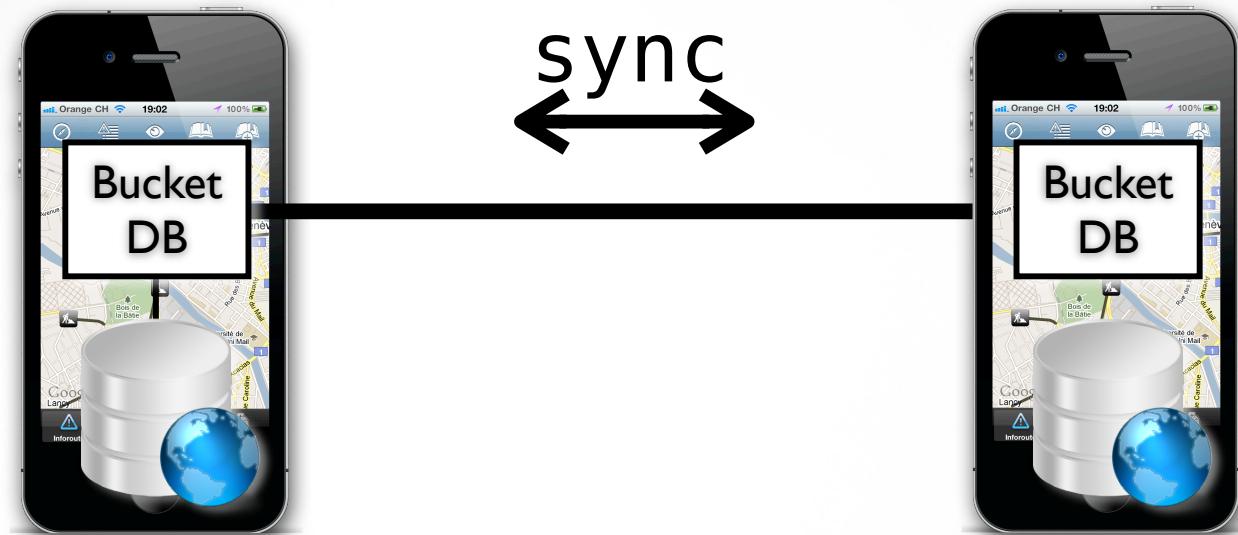
Content Distribution

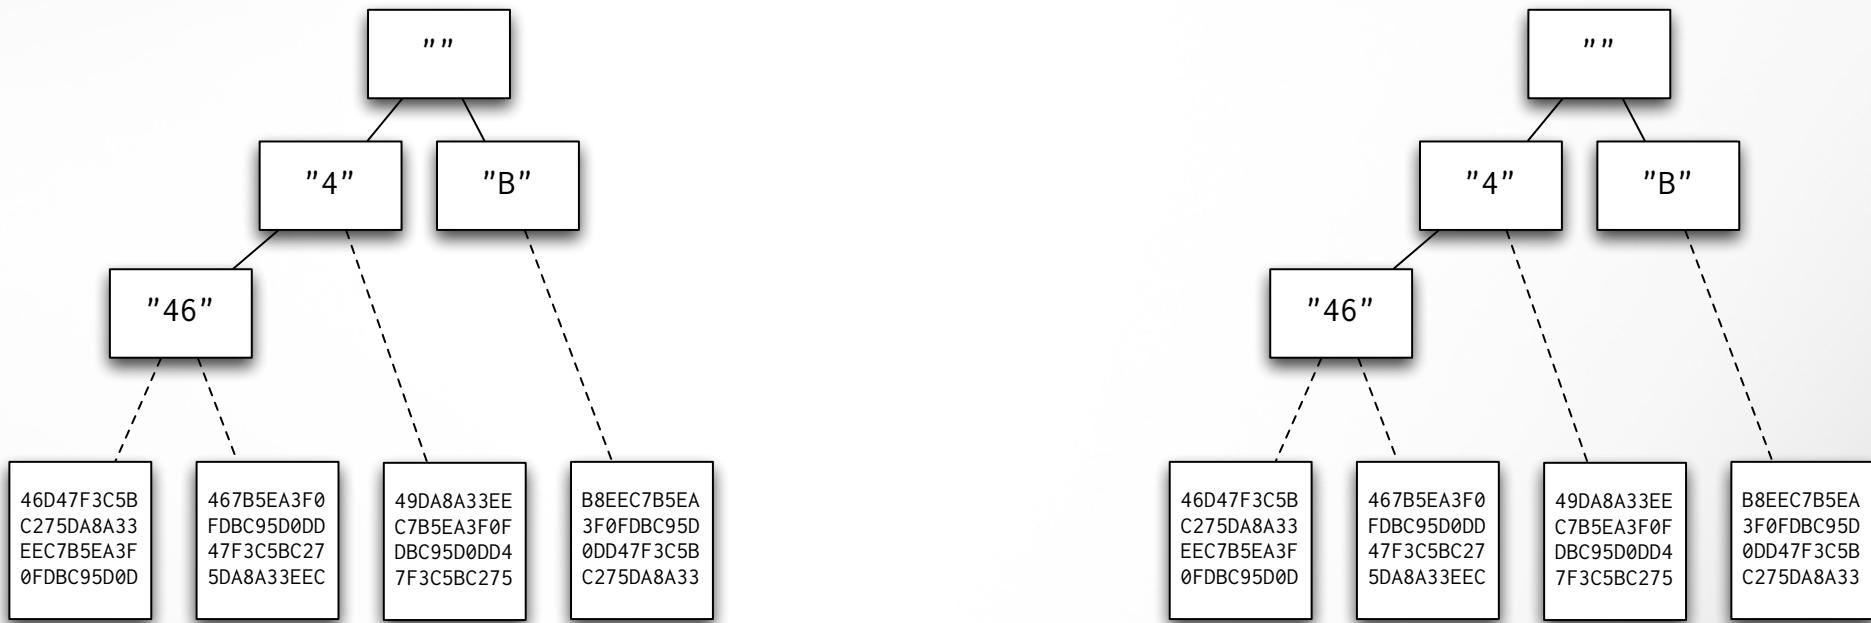


Reliable Queueing

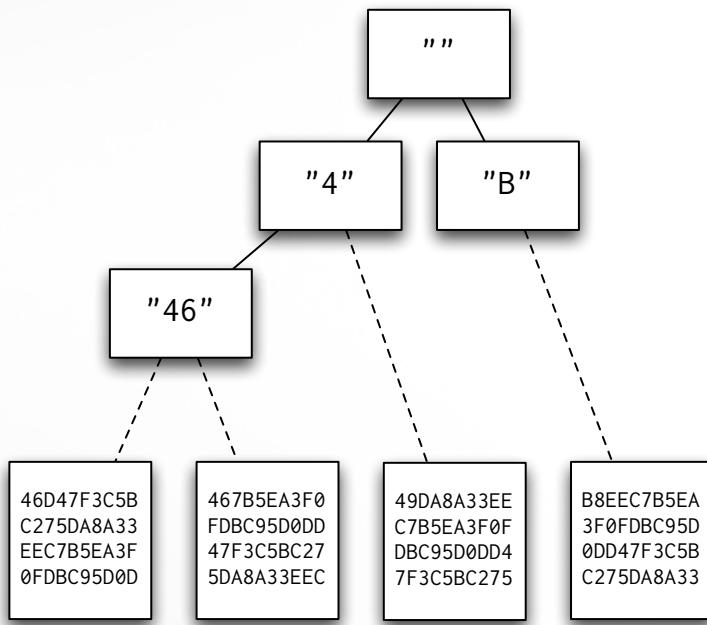


Peer Sync





BucketDB



- Client stores data in local data store as a HASH TREE
- Always “ready to sync” (no session state, no cpu/memory intensive computation)

DATA

ROW_ID	KEY	CONTENT TYPE	DATA

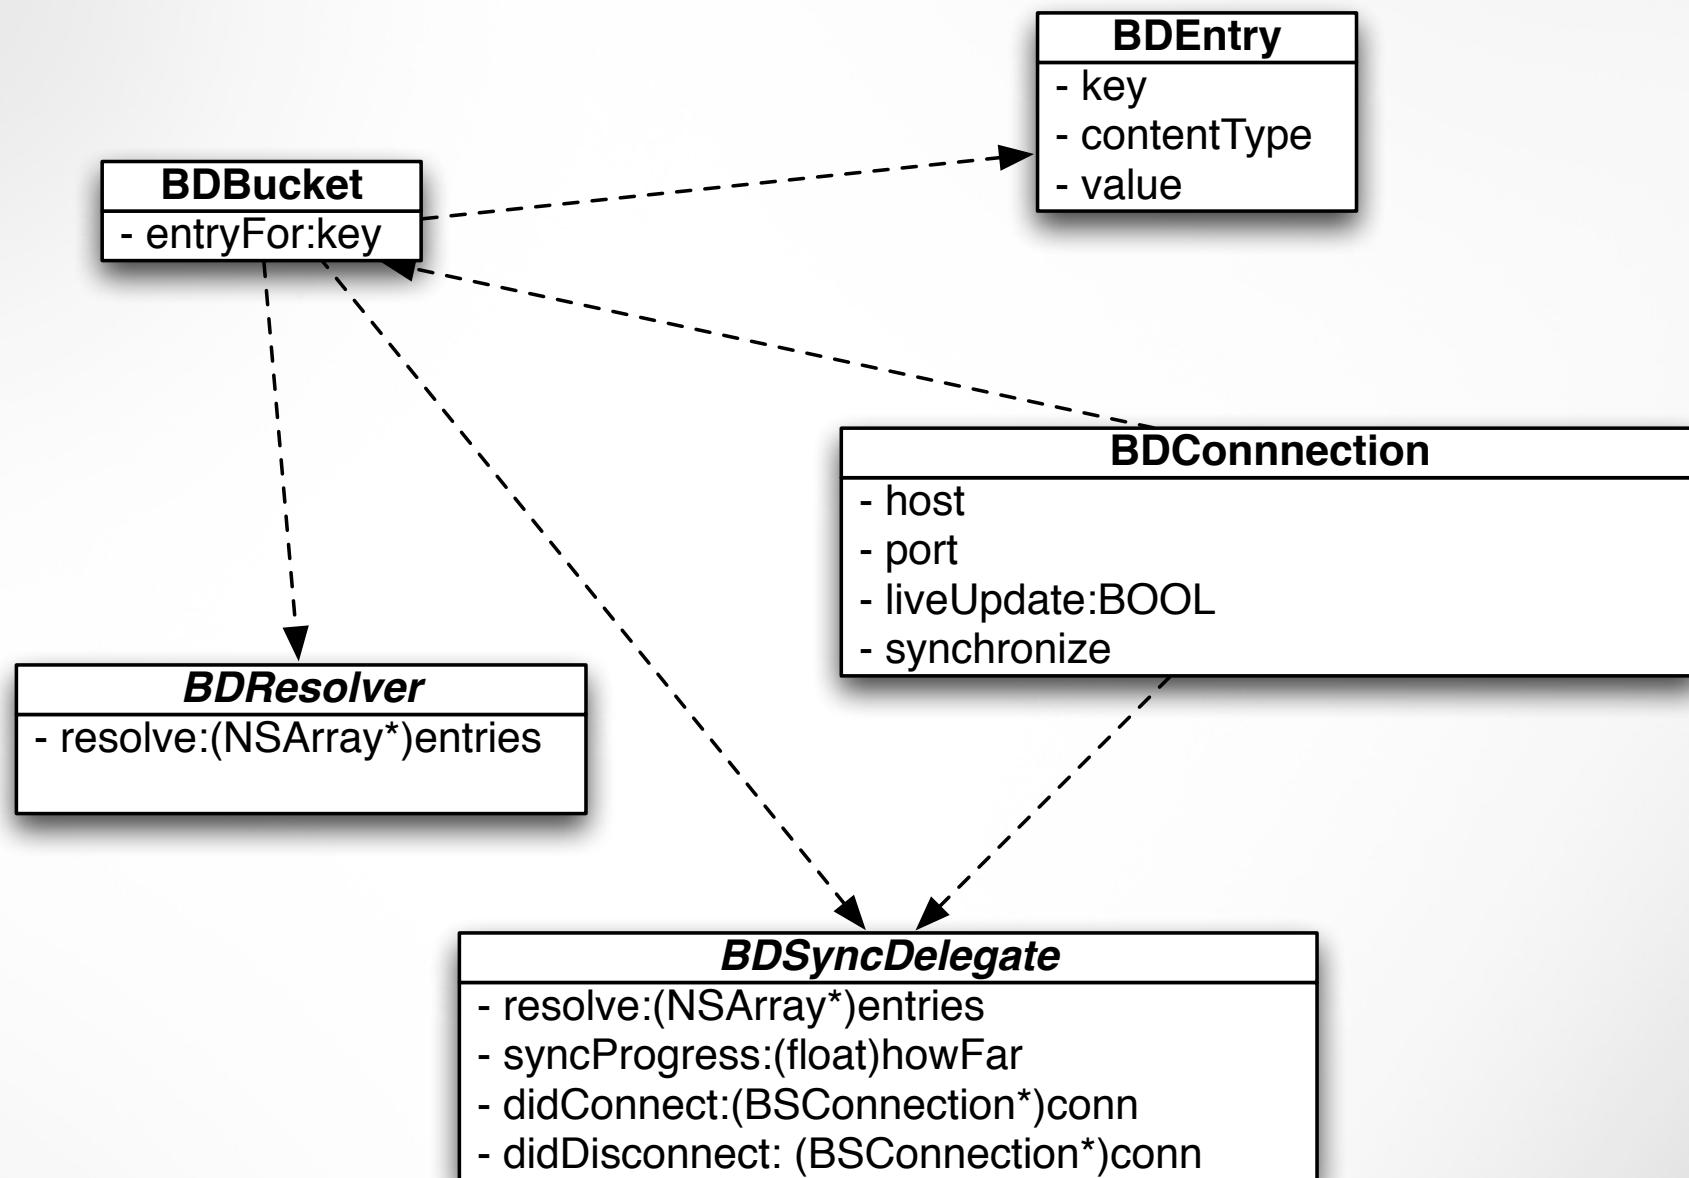
TRIGGER

MERKLE_LEAF

MERKLE_INNER

HKEY	VCLOCK	CHASH	KEY

PATH	CHASH	CHILDREN



Other Items to Discuss

- Security models
 - Provide own “socket factory” / credentials
 - Document- or Record-oriented
- Handling Deleted data
- Future: Index/Search, local M/R

Summary

- Riak Data Model
- RiakSync, a protocol for Key/Value synchronization
- BucketDB, Riak clients for mobile

Thank You

@drkrab



Thank You

@drkrab

